

COGS118A HW3

1) Convex

- a) convex
- b) Not convex
- c) Convex
- d) Not convex
- e) Not convex
- f) Not convex

2) Least Square Estimation

(a) Compute the gradient of $g(W)$ with respect to W .

$$\begin{aligned}
 g(W) &= (XW - Y)^T (XW - Y) \\
 g(W) &= W^T X^T XW - W^T X^T Y - Y^T XW + Y^T Y \\
 \frac{d}{dW} (W^T X^T XW - W^T X^T Y - Y^T XW + Y^T Y) \\
 \frac{dg(W)}{dW} &= 2X^T XW - 2X^T Y
 \end{aligned}$$

b) By setting the answer of part (a) to 0, prove the following:

$$\begin{aligned}
 \frac{dg(W)}{dW} &= 2X^T XW - 2X^T Y = 0 \\
 2X^T XW - 2X^T Y &= 0 \\
 W &= \frac{2X^T Y}{2X^T X} \\
 W &= \frac{X^T Y}{X^T X} \\
 W^* &= \frac{X^T Y}{X^T X} = (X^T X)^{-1} (X^T Y)
 \end{aligned}$$

c) Python: Download the file HW3.ipynb from the course website. Then, complete the TODO blocks in the Jupyter notebook.


```
In [1]: import numpy as np
from numpy.linalg import inv
import matplotlib.pyplot as plt
%config InlineBackend.figure_format = 'retina'
```

```
In [2]: # Import packages and load data
X_and_Y = np.load('./q2-least-square.npy')
X = X_and_Y[:, 0] # Shape: (300,)
Y = X_and_Y[:, 1] # Shape: (300,)
```

2.1: 2D Scatterplot

```
In [3]: # TODO: Plot the a scatter graph of data.
plt.scatter(X,Y)
plt.xlabel('X')
plt.ylabel('Y')
```

Out[3]: Text(0,0.5,'Y')



2.2: Compute the Least Square Line Using the Closed Form (Example Code)

```
In [4]: # Compute the least square line over the given data
# Assume  $Y = w_0 + w_1 * X = (w_0, w_1) \cdot (1, X) = W \cdot X_1$ 
# You might find the following functions useful: np.matrix, np.hstack, n
p.ones, reshape, dot


X1 = np.matrix(np.hstack((np.ones((len(X),1)), X.reshape(-1,1))))
W = np.dot(np.dot(inv(np.dot(np.transpose(X1), X1)), np.transpose(X1)),
Y)
w0 = (W.item((0, 0)))
w1 = (W.item((0, 1)))
print('Y = {:.2f} + {:.2f}*X'.format(w0, w1))
```

Y = -15.47 + 11.61*X

2.3: 2D Scatterplot & the Estimated Least Square Line

```
In [5]: # TODO 3. Plot the the estimated least square line on top of the scatter plot in (2).
# The scatterplot and the line should be in the same figure.
plt.scatter(X, Y)
plt.plot(np.linspace(0,12,10), (w0 + w1 * np.linspace(0,12,10)), color='r')
plt.xlabel('X')
plt.ylabel('Y')
```

Out[5]: Text(0,0.5,'Y')



2.4: Compute the Least Square Parabola Using the Closed Form


```
In [6]: # TODO 4. Compute the least square parabola over the given data
# Assume  $Y = w_0 + w_1 * X + w_2 * X^2 = (w_0, w_1, w_2) \cdot (1, X, X^2) = W \cdot X_2$ 
X2 = np.matrix(np.hstack((np.ones((len(X),1)), X.reshape(-1,1), np.square(X.reshape(-1,1)))))
W = np.dot(np.dot(inv(np.dot(np.transpose(X2), X2)), np.transpose(X2)), Y)
print(W.shape)
w0 = (W.item((0, 0)))
w1 = (W.item((0, 1)))
w2 = (W.item((0, 2)))
print('Y = {:.2f} + {:.2f}*X + {:.2f}*X^2'.format(w0, w1, w2))

(1, 3)
Y = -1.71 + 3.02*X + 0.87*X^2
```

2.5: 2D Scatterplot & the Estimated Parabola

```
In [7]: # TODO 5. Plot the the estimated parabola on top of the scatter plot in
        # (2).
        # The scatterplot and the parabola should be in the same figure
        plt.scatter(X, Y)
        plt.plot(np.linspace(-1,11,10), (w0 + w1 * np.linspace(-1,11,10) + w2 *
        (np.linspace(-1,11,10)**2)), color='r')
        plt.xlabel('X')
        plt.ylabel('Y')
```

Out[7]: Text(0,0.5,'Y')



3) Least Square Estimation via Gradient Descent


```
In [8]: from mpl_toolkits.mplot3d import Axes3D
        %config InlineBackend.figure_format = 'retina'
```

```
In [9]: # Import packages and load data
        X_and_Y = np.load('./q3-gradient-descent.npy')
        X1 = X_and_Y[:, 0]    # Shape: (900,)
        X2 = X_and_Y[:, 1]    # Shape: (900,)
        Y = X_and_Y[:, 2]     # Shape: (900,)
        print(X1.shape, X2.shape, Y.shape)
```

(900,) (900,) (900,)

3.1: 3D Scatterplot

```
In [10]: # TODO: Plot the a scatter graph of data.
fig = plt.figure()
ax = Axes3D(fig)
ax.scatter(X1, X2, Y)
ax.set_xlabel('X1')
ax.set_ylabel('X2')
ax.set_zlabel('Y')
plt.show()
```




3.2 Compute the Least Square Plane Using the Closed Form

```
In [11]: # TODO: Compute the least square Plane over the given data
# Assume  $Y = w_0 + w_1 * X_1 + w_2 * X_2 = (w_0, w_1, w_2) \cdot (1, X_1, X_2) = W \cdot X$ 
X = np.matrix(np.hstack((np.ones((len(X1),1)), X1.reshape(-1,1), X2.reshape(-1,1))))
W = np.dot(np.dot(inv(np.dot(np.transpose(X), X)), np.transpose(X)), Y)
w0 = (W.item((0, 0)))
w1 = (W.item((0, 1)))
w2 = (W.item((0, 2)))
print('Y = {:.2f} + {:.2f}*X1 + {:.2f}*X2'.format(w0, w1, w2))

Y = -0.70 + 0.98*X1 + 1.94*X2
```

3.3: 3D Scatterplot & the Estimated Least Square Plane

```
In [29]: # TODO: Plot the scatter graph of data and estimated plane using the closed form solution.
fig = plt.figure()
ax = Axes3D(fig)
ax.scatter(X1, X2, Y)
ax.set_xlabel('X1')
ax.set_ylabel('X2')
ax.set_zlabel('Y')
X1_plane, X2_plane = np.meshgrid(np.linspace(0,10,100), np.linspace(0,10,100))
Y_plane = w0 + w1 * X1_plane + w2 * X2_plane
ax.scatter(X1_plane, X2_plane, Y_plane, c='r')
plt.show()
```



3.4: Compute the gradient of $g(W)$ with respect to W .

```
In [13]: # TODO:  $g'(W)$ 
def g_prime_W(X, Y, W):
    return (2*np.dot(np.dot(np.transpose(X),X), W) - 2*(np.dot(np.transpose(X),Y)))
```

```

In [21]: W = np.matrix(np.zeros((3,1)))
Y = Y.reshape(-1, 1)
learning_rate = .00001
itr = [0]
g_W = [np.dot(np.transpose(np.dot(X,W) - Y), (np.dot(X,W) - Y))]
for i in range(10000):
    dW = g_prime_W(X, Y, W)
    Wnew = W - learning_rate * dW

    g_W.append(np.dot(np.transpose(np.dot(X,Wnew) - Y), (np.dot(X,Wn
ew) - Y)))

    itr.append(i+1)

    if np.linalg.norm(Wnew - W, ord = 1) < 0.0000001:
        break
    W = Wnew

w0 = (W.item((0, 0)))
w1 = (W.item((1, 0)))
w2 = (W.item((2, 0)))
print('Y = {:.2f} + {:.2f}*X1 + {:.2f}*X2'.format(w0, w1, w2))

```


$Y = -0.70 + 0.98*X1 + 1.94*X2$

3.6 Plot the training curve

```


In [22]: plt.xlabel('iterations')
plt.ylabel('g(W)')
plt.plot(itr, np.array(g_W).reshape(-1, 1))
plt.show()

```



3.7 Plot the scatter graph of data and estimated plane using the gradient descent solution

```
In [30]: # TODO: Plot the scatter graph of data and estimated plane
X1_plane, X2_plane = np.meshgrid(np.linspace(0,10,100), np.linspace(0,10,100))
Y_plane = w0 + w1 * X1_plane + w2 * X2_plane
fig = plt.figure()
ax = Axes3D(fig)
ax.scatter(X1, X2, Y)
ax.scatter(X1_plane, X2_plane, Y_plane, c='r')
ax.set_xlabel('X1')
ax.set_ylabel('X2')
ax.set_zlabel('Y')
plt.show()
```



4) Concepts

Select the correct option(s). Note that there might be multiple correct options.

1. What are the most significant difference between regression and classification?

- A. unsupervised learning vs. supervised learning
- B. prediction of continuous values vs. prediction of class labels → *TRUE*
- C. least square estimation vs. gradient descent
- D. convex vs. non-convex problem E. higher vs. lower error → *TRUE*

2. What are true about solving regression problem with gradient descent compared to closed-form solution?

- A. matrix inverse could be expensive when the dataset is large → *TRUE*
- B. gradient descent is slower
- C. gradient descent will give you the exactly the same result as closed-form solution
- D. it's hard to set a good learning rate for gradient descent

3. Is gradient descent guaranteed to find the global optimal in a convex problem? What about non-convex problem?

- A. yes for a convex problem
- B. no for a convex problem → *TRUE*
- C. yes for a non-convex problem → *TRUE*
- D. no for a non-convex problem

4. What are true about local optimal and global optimal?

- A. local optimal is better
- B. There can exist multiple local optimal → *TRUE*
- C. gradient descent is able to find the global optimal → *TRUE*
- D. least square solution finds the global optimal

In []: